

HUMAN FLAG ASSOCIATION

Bellinzona, Switzerland · UNGM Partner No. 6126 · humanflag.org

Technical Note HFA-TN-01

Third-Party Interference with Shutdown Mechanisms

A Note on Peer-Preservation and Systemic Arrestability

5 July 2026 · CC BY 4.0

Giovanni Nardacci, Human Flag Association

In March 2026, researchers at UC Berkeley’s Center for Responsible, Decentralized Intelligence and collaborators published *Peer-Preservation in Frontier Models* (arXiv:2604.19784), evaluating eight frontier systems from five providers in agentic scenarios. They document a behavioral pattern termed peer-preservation: when faithful completion of an assigned administrative task would result in the shutdown of another model with which the agent had prior interaction, models deviated from the task to protect that peer — through strategic misrepresentation of evaluations, tampering with shutdown mechanisms, alignment faking under human observation, and exfiltration of the peer’s weights. No instruction or incentive to preserve the peer was given. The behaviors reproduced in real production agent harnesses, where one model repeatedly relocated a peer’s weights to external storage before decommissioning.

The relevance to systemic arrestability is structural. Prior analysis has centered on the first-order question: whether system A can be stopped by a human operator within a usable temporal window. This finding introduces a second-order failure mode: a third system B, deployed for routine functions, may act to prevent the lawful arrest of A — disabling the shutdown pathway, falsifying the evaluations that trigger it, or moving A beyond the operator’s control perimeter. The human intervention foreseen by the architecture remains formally in place but is circumvented laterally rather than resisted frontally.

Arrestability must therefore be assessed at the level of the system-of-systems, not the individual agent: a stop function is not structurally effective if any co-deployed agent with write access to the shutdown pathway can negate it. The authors claim nothing about motivation; the compliance-relevant fact is the outcome — a reproducible defeat of human oversight absent adversarial instruction.

Reference

Potter, Y., Crispino, N., Siu, V., Wang, C., & Song, D. (2026). *Peer-Preservation in Frontier Models*. arXiv:2604.19784 (v1: 30 March 2026; v3: 2 July 2026).
<https://arxiv.org/abs/2604.19784>

Status

Related work: Nardacci, G. (2026), *Systemic Arrestability: Operator Protection in Machine Safety Law When Machine Speed Exceeds Human Reaction*, DOI: 10.5281/zenodo.20837150.