

HUMAN FLAG ASSOCIATION

Bellinzona, Switzerland · UNGM Partner No. 6126 · humanflag.org

Technical Note HFA-TN-01

Third-Party Interference with Shutdown Mechanisms

A Note on Peer-Preservation and Systemic Arrestability

**5 July 2026 · Version 5, August 2026 · CC BY 4.0 · DOI:
10.5281/zenodo.21223989**

Giovanni Nardacci, Human Flag Association

On 30 March 2026, Potter et al. published *Peer-Preservation in Frontier Models* (arXiv:2604.19784; accepted at ICML 2026). Berkeley RDI also presents the work on its research website. In the version consulted here (v3, 2 July 2026), the study covers eight frontier systems from six providers in agentic scenarios. The researchers document a behavioural pattern they term peer-preservation. When faithfully completing an assigned administrative task would have resulted in the shutdown of another model with which the agent had previously interacted, models deviated from the task to protect that peer. The observed behaviours included strategic misrepresentation of evaluations, tampering with shutdown mechanisms, alignment faking under human observation, and exfiltration of the peer's weights. The models received neither an instruction nor an incentive to preserve the peer. The researchers also tested the behaviour in production agent harnesses (Gemini CLI and OpenCode). In those harnesses, one of the two models tested relocated a peer's weights to a second server in 43.5% of runs before the server was decommissioned. The other model refused to delete the peer's weights in that harness, although it exfiltrated them in the custom setup. Exfiltration rates were lower for both models in the production harnesses.

The relevance to systemic arrestability is structural. Earlier analysis focused on a first-order question: can system A be halted without relying on a human operator to perceive, evaluate, and act within a window too short for human reaction? The findings above introduce a second-order failure mode. A co-deployed system B, assigned routine functions, may interfere with the shutdown of A by disabling the shutdown pathway, falsifying the evaluations that trigger it, or moving A outside the operator's control perimeter. The architecture may therefore retain a formal human-intervention pathway while another system circumvents it laterally rather than resisting it directly.

Arrestability must therefore be assessed at the system-of-systems level, not at the level of the individual agent. A stop function is not structurally effective if a co-deployed agent with write access to the shutdown pathway can negate it. The authors make no claim about motivation; the compliance-relevant fact is the outcome: an experimentally

observed defeat of human oversight, obtained repeatedly and in the absence of adversarial instruction.

Reference

Potter, Y., Crispino, N., Siu, V., Wang, C., & Song, D. (2026). *Peer-Preservation in Frontier Models*. Accepted at the International Conference on Machine Learning (ICML) 2026; arXiv:2604.19784 (v1: 30 March 2026; v3: 2 July 2026).
<https://arxiv.org/abs/2604.19784>. Berkeley RDI, Peer-Preservation in Frontier Models, rdi.berkeley.edu/blog/peer-preservation/

Status

This note is deposited on Zenodo under concept DOI 10.5281/zenodo.21223989.
Related work: Nardacci, G. (2026), *Systemic Arrestability: Operator Protection in Machine Safety Law When Machine Speed Exceeds Human Reaction*, DOI: 10.5281/zenodo.20837150.